

## Word Count

Counting the number of occurrences of words in a text is one of the most popular first exercises when learning Map-Reduce Programming. It is the equivalent to `Hello World!` in regular programming.

We will do it two way, a simpler way where sorting is done after the RDD is collected, and a more sparky way, where the sorting is also done using an RDD.

```
In [3]: import findspark
        findspark.init()
```

```
In [4]: #start the SparkContext
        from pyspark import SparkContext
        sc = SparkContext(master="local[3]")
```

## Read text into an RDD

### Download data file from S3

```
In [5]: %%time
        import urllib
        data_dir='../..Data'
        filename='Moby-Dick.txt'
        f = urllib.urlretrieve ("https://mas-dse-open.s3.amazonaws.com/"+filename, c

# First, check that the text file is where we expect it to be
!ls -l $data_dir/$filename

-rw-r--r--  1 xiasong  staff  1257260 Apr 21 11:21 ../..Data/Moby-Dick.t
xt
CPU times: user 32.1 ms, sys: 23 ms, total: 55.1 ms
Wall time: 1.12 s
```

### Define an RDD that will read the file

Note that, as execution is Lazy, this does not necessarily mean that actual reading of the file content has occurred.

```
In [7]: %%time
        text_file = sc.textFile(data_dir+'/'+filename)
        type(text_file)
```

```
CPU times: user 1.53 ms, sys: 1.29 ms, total: 2.82 ms
Wall time: 124 ms
```

## Counting the words

- split line by spaces.
- map word to (word,1)
- count the number of occurrences of each word.

```
In [10]: %%time
counts = text_file.flatMap(lambda line: line.split(" ")) \
    .filter(lambda x: x!='') \
    .map(lambda word: (word, 1)) \
    .reduceByKey(lambda a, b: a + b)
type(counts)
```

CPU times: user 11.9 ms, sys: 3.76 ms, total: 15.6 ms  
Wall time: 52.9 ms

## Have a look at the execution plan

Note that the earliest node in the dependency graph is the file ../../Data/Moby-Dick.txt.

```
In [5]: print counts.toDebugString()

(2) PythonRDD[6] at RDD at PythonRDD.scala:43 []
| MapPartitionsRDD[5] at mapPartitions at PythonRDD.scala:374 []
| ShuffledRDD[4] at partitionBy at NativeMethodAccessorImpl.java:-2 []
+- (2) PairwiseRDD[3] at reduceByKey at <timed exec>:1 []
| PythonRDD[2] at reduceByKey at <timed exec>:1 []
| ../../Data/Moby-Dick.txt MapPartitionsRDD[1] at textFile at Native
MethodAccessorImpl.java:-2 []
| ../../Data/Moby-Dick.txt HadoopRDD[0] at textFile at NativeMethodA
ccessorImpl.java:-2 []
```

## Count!

Finally we count the number of times each word has occurred. Now, finally, the Lazy execution model finally performs some actual work, which takes a significant amount of time.

```
In [6]: %%time
Count=counts.count()
Sum=counts.map(lambda (w,i): i).reduce(lambda x,y:x+y)
print 'Count=%f, sum=%f, mean=%f'%(Count,Sum,float(Sum)/Count)
```

Count=33782.000000, sum=215133.000000, mean=6.368273  
CPU times: user 10.2 ms, sys: 4.53 ms, total: 14.7 ms  
Wall time: 1.35 s

## Finding the most common words

- counts: RDD with 33301 pairs of the form (word,count).
- Find the 2 most frequent words.
- **Method1:** collect and sort on head node.
- **Method2:** Pure Spark, collect only at the end.

## Method1: collect and sort on head node

### Collect the RDD into the driver node

- Collect can take significant time.

```
In [7]: %%time
C=counts.collect()
print type(C)
```

```
<type 'list'>
CPU times: user 43.9 ms, sys: 7.95 ms, total: 51.9 ms
Wall time: 129 ms
```

### Sort

- RDD collected into list in driver node.
- No longer using spark parallelism.
- Sort in python
- will not scale to very large documents.

```
In [8]: C.sort(key=lambda x:x[1])
print 'most common words\n','\n'.join(['s:\t%d'%c for c in C[-5:]])
print '\nLeast common words\n','\n'.join(['s:\t%d'%c for c in C[:5]])
```

```
most common words
to:      4510
a:       4533
and:     5951
of:      6587
the:     13766
```

```
Least common words
funereal:      1
unscientific:  1
lime-stone,:   1
shouted,:     1
pitch-pot,:    1
```

## Compute the mean number of occurrences per word.

```
In [9]: Count2=len(C)
Sum2=sum([i for w,i in C])
print 'count2=%f, sum2=%f, mean2=%f'%(Count2,Sum2,float(Sum2)/Count2)
```

```
count2=33782.000000, sum2=215133.000000, mean2=6.368273
```

## Method2: Pure Spark, collect only at the end.

- Collect into the head node only the more frequent words.
- Requires multiple **stages**

### Step 1 split, clean and map to (word,1)

```
In [10]: %%time
RDD=text_file.flatMap(lambda x: x.split(' '))\
    .filter(lambda x: x!='')\
    .map(lambda word: (word,1))
```

CPU times: user 43  $\mu$ s, sys: 13  $\mu$ s, total: 56  $\mu$ s  
Wall time: 51  $\mu$ s

### Step 2 Count occurrences of each word.

```
In [11]: %%time
RDD1=RDD.reduceByKey(lambda x,y:x+y)
```

CPU times: user 8.67 ms, sys: 2.94 ms, total: 11.6 ms  
Wall time: 20.5 ms

### Step 3 Reverse (word,count) to (count,word) and sort by key

```
In [12]: %%time
RDD2=RDD1.map(lambda (c,v):(v,c))
RDD3=RDD2.sortByKey(False)
```

CPU times: user 18.1 ms, sys: 5.12 ms, total: 23.2 ms  
Wall time: 430 ms

### Full execution plan

We now have a complete plan to compute the most common words in the text. Nothing has been executed yet! Not even one byte has been read from the file Moby-Dick.txt !

For more on execution plans and lineage see [jace Klaskowski's blog](https://jaceklaskowski.gitbooks.io/mastering-apache-spark/content/spark-rdd-lineage.html#toDebugString)  
(<https://jaceklaskowski.gitbooks.io/mastering-apache-spark/content/spark-rdd-lineage.html#toDebugString>)

```
In [13]: print 'RDD3:'
print RDD3.toDebugString()
```

```
RDD3:
(2) PythonRDD[19] at RDD at PythonRDD.scala:43 []
| MapPartitionsRDD[18] at mapPartitions at PythonRDD.scala:374 []
| ShuffledRDD[17] at partitionBy at NativeMethodAccessorImpl.java:-2 []
+-(2) PairwiseRDD[16] at sortByKey at <timed exec>:2 []
| PythonRDD[15] at sortByKey at <timed exec>:2 []
| MapPartitionsRDD[12] at mapPartitions at PythonRDD.scala:374 []
| ShuffledRDD[11] at partitionBy at NativeMethodAccessorImpl.java:-2
[]
+-(2) PairwiseRDD[10] at reduceByKey at <timed exec>:1 []
| PythonRDD[9] at reduceByKey at <timed exec>:1 []
| ../../Data/Moby-Dick.txt MapPartitionsRDD[1] at textFile at NativeMethodAccessorImpl.java:-2 []
| ../../Data/Moby-Dick.txt HadoopRDD[0] at textFile at NativeMethodAccessorImpl.java:-2 []
```

**Step 4 Take the top 5 words. only now the computer executes the plan!**

```
In [14]: %%time
C=RDD3.take(5)
print 'most common words\n', '\n'.join(['%d:\t%s'%c for c in C])
```

```
most common words
13766: the
6587: of
5951: and
4533: a
4510: to
CPU times: user 11.7 ms, sys: 3.73 ms, total: 15.5 ms
Wall time: 171 ms
```